

# Sampling Distributions

Eli Gurarie

StatR 101 - Lecture 8b  
November 15, 2012

November 15, 2012

PROFESSIONAL & CONTINUING EDUCATION

UNIVERSITY *of* WASHINGTON




# Review of Definitions

- a **parameter** is a number that describes a “true” distribution
  - $\mu$  - mean,  $\sigma$  - standard deviation
  - $\alpha$ ,  $\beta$ ,  $\gamma$ , etc. in continuous distributions
  - $p$  - probability of a Bernoulli or Binomial
- a **statistic** is *any numerical summary of data*
  - $\bar{X}$  - sample mean
  - $s_x$  - sample standard deviation
  - $X_{min}$ ,  $X_{max}$

# Review of Definitions

## Data



## Statistics

sample mean:  $\bar{x} = 112.5$


sample variance:  $s_x^2 = 29.8$

sample s. d.:  $s_x = 5.46$



# Review of Definitions

## Data




## Statistics

sample mean:  $\bar{x} = 112.5$

sample variance:  $s_x^2 = 29.8$

sample s. d.:  $s_x = 5.46$

## Model



## Parameters

theoretical mean:  $\mu = 112.5$

theoretical variance:  $\sigma^2 = 29.8$

theoretical s.d.:  $\sigma = 5.46$

# Inference ...

is when you try to say something about a **population** when all you have data on is a **sample**.

- We have a population of interest
- The population has some *true* structure - e.g. a *true* distribution with a *true* parameters (e.g. mean  $\mu$ , standard deviation  $\sigma$ ,  $p$  etc.).
- We collect a *random sample* from the population ( $X_1, X_2, X_3 \dots X_n$ ) - because we can't measure the entire population - and calculate a *statistic*,
- We determine how *good* the statistic is at **estimating** the parameter
  - How **accurate** is it? - i.e. how **biased**
  - How **precise** is it? - i.e. how big is the **margin of error** or **confidence interval**.

# Inference ...

is when you try to say something about a **population** when all you have data on is a **sample**.

- We have a population of interest
- The population has some *true* structure - e.g. a *true* distribution with a *true* parameters (e.g. mean  $\mu$ , standard deviation  $\sigma$ ,  $p$  etc.).
- We collect a *random sample* from the population ( $X_1, X_2, X_3 \dots X_n$ ) - because we can't measure the entire population - and calculate a *statistic*,
- We determine how *good* the statistic is at **estimating** the parameter
  - How **accurate** is it? - i.e. how **biased**
  - How **precise** is it? - i.e. how big is the **margin of error** or **confidence interval**.

# Inference ...

is when you try to say something about a **population** when all you have data on is a **sample**.

- We have a population of interest
- The population has some *true* structure - e.g. a *true* distribution with a *true* parameters (e.g. mean  $\mu$ , standard deviation  $\sigma$ ,  $p$  etc.).
- We collect a *random sample* from the population ( $X_1, X_2, X_3 \dots X_n$ ) - because we can't measure the entire population - and calculate a *statistic*,
- We determine how *good* the statistic is at **estimating** the parameter
  - How **accurate** is it? - i.e. how **biased**
  - How **precise** is it? - i.e. how big is the **margin of error** or **confidence interval**.

# Inference ...

is when you try to say something about a **population** when all you have data on is a **sample**.

- We have a population of interest
- The population has some *true* structure - e.g. a *true* distribution with a *true* parameters (e.g. mean  $\mu$ , standard deviation  $\sigma$ ,  $p$  etc.).
- We collect a *random sample* from the population ( $X_1, X_2, X_3 \dots X_n$ ) - because we can't measure the entire population - and calculate a *statistic*,
- We determine how *good* the statistic is at **estimating** the parameter
  - How **accurate** is it? - i.e. how **biased**
  - How **precise** is it? - i.e. how big is the **margin of error** or **confidence interval**.



The most common parameter of interest is ...

$$\mu$$

The best **estimator** of the parameter is ...

$$\overline{X} \approx \mu$$

The big question is ...

How good is  $\bar{X}$  at estimating  $\mu$ ?

## The reasoning

- The values  $X_1, X_2, X_3 \dots X_n$  are all *random, independent* observations from some distribution  $X \sim f(x)$  with some mean value  $\mu$ . (note that  $X$  can be *any* distribution).
- Recall that:

$$\bar{X} = \frac{1}{n} \sum X_i$$

- So:  $\bar{X}$  is also a random variable.

## Expectation of $\overline{X}$

- Recall the rules of summing and multiplying expectations:

$$E(X + Y) = E(X) + E(Y)$$

$$E(aX) = aE(X)$$

$$E(a + bX) = a + bE(X)$$

- So:

$$\begin{aligned} E(\overline{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) \\ &= \frac{1}{n} \sum_{i=1}^n \mu = \frac{1}{n} n\mu = \mu \end{aligned}$$

## Expectation of $\overline{X}$

- Recall the rules of summing and multiplying expectations:

$$E(X + Y) = E(X) + E(Y)$$

$$E(aX) = aE(X)$$

$$E(a + bX) = a + bE(X)$$

- So:

$$\begin{aligned} E(\overline{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) \\ &= \frac{1}{n} \sum_{i=1}^n \mu = \frac{1}{n} n\mu = \mu \end{aligned}$$

## Expectation of $\bar{X}$

So (using expectation arithmetic):

$$E(\bar{X}) = \mu$$

Therefore, we say that:  $\bar{X}$  is an **unbiased estimator** of  $\mu$ , because its *expectation* is exactly the parameter that we are estimating.

## Variance of $\bar{X}$

- Here are the complete rules of variance arithmetic:

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(XY)$$

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y) - 2\text{Cov}(XY)$$

$$\text{Var}(a + X) = \text{Var}(X)$$

$$\text{Var}(bX) = b^2\text{Var}(X)$$

$$\text{Var}(aX + bY) = a^2\text{Var}(X) + b^2\text{Var}(Y) + 2ab\text{Cov}(X, Y)$$



## Brief aside on $\text{Cov}(X, Y)$

The **covariance** of two random variables  $X$  and  $Y$  is given by:

$$\begin{aligned}\text{Cov}(X, Y) &= E[(X - E(X))(Y - E(Y))] \\ &= E(XY) - E(X)E(Y)\end{aligned}$$

For now, we will not worry about covariances at all, because we will assume that the observations we make are independent.

# Rules of expectations and variances (under independence)

- Rules of expectation arithmetic:

$$E(X + Y) = E(X) + E(Y)$$

$$E(aX) = aE(X)$$

$$E(a + bX) = a + bE(X)$$

- Rules of variance arithmetic:

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$$

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y)$$

$$\text{Var}(a + X) = \text{Var}(X)$$

$$\text{Var}(bX) = b^2\text{Var}(X)$$

$$\text{Var}(aX + bY) = a^2\text{Var}(X) + b^2\text{Var}(Y)$$

## Back to the variance of $\bar{X}$

- Recall relevant rules of variance arithmetic:

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(XY)$$

$$\text{Var}(bX) = b^2\text{Var}(X)$$

- So:

$$\begin{aligned}\text{Var}(\bar{X}) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\&= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\&= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \\&= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{1}{n^2} n\sigma^2 = \frac{\sigma^2}{n}\end{aligned}$$

## Back to the variance of $\bar{X}$

- Recall relevant rules of variance arithmetic:

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(XY)$$

$$\text{Var}(bX) = b^2\text{Var}(X)$$

- So:

$$\begin{aligned}\text{Var}(\bar{X}) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\&= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\&= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \\&= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{1}{n^2} n\sigma^2 = \frac{\sigma^2}{n}\end{aligned}$$

## Expectation and variance of $\bar{X}$

If  $X$  is a population with true mean  $\mu$  and variance  $\sigma^2$ , and  $X_1, X_2, X_3 \dots X_n$  are observations of that population, and the sample mean  $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$  then:

- $E(\bar{X}) = \mu$
- $\text{Var}(\bar{X}) = \frac{\sigma^2}{n}$

Note that we derived these from the “arithmetic” rules of expectation and variance.

## The standard deviation of the sample mean $\sigma_{\bar{X}}$

- The precision of our estimate  $\bar{X}$  is determined by the size of our sample:

$$\sigma_{\bar{X}} = \frac{\sigma_X}{\sqrt{n}}$$

- The larger the sample  $n$ , the smaller  $\sigma_{\bar{X}}$ .
- So, because  $E(\bar{X}) = \mu$  it is an **accurate (unbiased)** estimator
- The **precision** of  $E(\bar{X})$  depends (as the inverse square root) on  $n$ .

## The standard deviation of the sample mean $\sigma_{\bar{X}}$

- The precision of our estimate  $\bar{X}$  is determined by the size of our sample:

$$\sigma_{\bar{X}} = \frac{\sigma_X}{\sqrt{n}}$$

- The larger the sample  $n$ , the smaller  $\sigma_{\bar{X}}$ .
- So, because  $E(\bar{X}) = \mu$  it is an **accurate (unbiased)** estimator
- The **precision** of  $E(\bar{X})$  depends (as the inverse square root) on  $n$ .

## Recall the central limit theorem:

### Central Limit Theorem (CLT)

If  $X_1, X_2, X_3 \dots X_n$  are **any, independent, identically distributed (iid)** random variables with mean  $\mu_x$  and standard deviation  $\sigma_x$ , and

$$Y = \sum_{i=1}^n X_i$$

then, as  $n$  becomes large

$$Y \sim \mathcal{N}(n\mu_x, n\sigma_x^2)$$

In words: The sum of random variables approximates a normal distribution no matter what the variable is.



## A stronger statement about $\bar{X}$

The asymptotic distribution of  $\bar{X}$

The distribution of  $\bar{X}$  is *approximately normal* with:



- $\bar{X} \sim \mathcal{N}(\mu, \frac{\sigma^2}{n})$

This is a direct consequence of the central limit theorem.

## In summary

- Arithmetic rules and central limit theorem let us say anything about the **sampling distribution** of  $\bar{X}$ .
- So we can solve any probability problem related to  $\bar{X}$ .

## Example with melatonin



### Melatonin


From Wikipedia, the free encyclopedia

*Not to be confused with [Melanin](#) or [Melanotan](#).*

**Melatonin** <sup>i</sup>/ˌmɛləˈtoʊnɪn/, also known chemically as ***N*-acetyl-5-methoxytryptamine**,<sup>[1]</sup> is a naturally occurring compound found in animals, plants and microbes.<sup>[2][3]</sup> In animals, circulating levels of the hormone melatonin vary in a daily cycle, thereby allowing the [entrainment](#) of the [circadian rhythms](#) of several biological functions.<sup>[4]</sup>

## Example with melatonin

Histogram of 40 times until sleep



## Example with melatonin

- Time to fall asleep for all humans:  $\mu = 15, \sigma = 10$
- If melatonin has no effect on time to sleep, then

$$X_1, X_2, \dots, X_{40} \sim \text{Some Distribution}(\mu = 15, \sigma = 10)$$

- But, by CLT

$$\begin{aligned}\bar{X} &\sim \mathcal{N}\left(\mu = 15, \sigma = \frac{10}{\sqrt{40}}\right) \\ &\mathcal{N}(\mu = 15, \sigma = 1.58)\end{aligned}$$

- So:

$$\begin{aligned}P(\bar{X} \leq 5.9) &= \text{pnorm}(5.9, \text{mean}=15, \text{sd}=1.58) \\ &= 4.218331\text{e-}09 \approx 0\end{aligned}$$

Another common parameter of interest is ...

$$\sigma^2$$

How good is the **estimator** ...

$$S^2 \approx \sigma^2$$

## What about $s^2$

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

Recall:

$$\sigma^2 = E((X - \mu)^2); \frac{\sigma^2}{n} = E((\bar{X} - E(\bar{X}))^2) = E((\bar{X} - \mu)^2)$$

$$\begin{aligned} E(s^2) &= E\left(\frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2\right) = \frac{1}{n-1} \sum_{i=1}^n E((X_i - \bar{X})^2) \\ &= \frac{1}{n-1} \sum_{i=1}^n E((X_i - \mu)^2 - (\mu - \bar{X})^2) \\ &= \frac{1}{n-1} \left( \sum_{i=1}^n E((X_i - \mu)^2) - \sum_{i=1}^n E((\bar{X} - \mu)^2) \right) \\ &= \frac{1}{n-1} \left( \sum_{i=1}^n \sigma^2 - \sum_{i=1}^n \frac{1}{n} \sigma^2 \right) \\ &= \frac{1}{n-1} (n\sigma^2 - \sigma^2) = \frac{n-1}{n-1} \sigma^2 = \sigma^2 \end{aligned}$$



## What about $s^2$

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

Recall:

$$\sigma^2 = E((X - \mu)^2); \frac{\sigma^2}{n} = E((\bar{X} - E(\bar{X}))^2) = E((\bar{X} - \mu)^2)$$

$$\begin{aligned} E(s^2) &= E\left(\frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2\right) = \frac{1}{n-1} \sum_{i=1}^n E((X_i - \bar{X})^2) \\ &= \frac{1}{n-1} \sum_{i=1}^n E((X_i - \mu)^2 - (\mu - \bar{X})^2) \\ &= \frac{1}{n-1} \left( \sum_{i=1}^n E((X_i - \mu)^2) - \sum_{i=1}^n E((\bar{X} - \mu)^2) \right) \\ &= \frac{1}{n-1} \left( \sum_{i=1}^n \sigma^2 - \sum_{i=1}^n \frac{1}{n} \sigma^2 \right) \\ &= \frac{1}{n-1} (n\sigma^2 - \sigma^2) = \frac{n-1}{n-1} \sigma^2 = \sigma^2 \end{aligned}$$

## What about $s^2$

So:

$$E(s^2) = \sigma^2$$

This means:  $s^2$  is an **unbiased estimator** of  $\sigma^2$ .

The  $n - 1$  in the denominator in  $s^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$  is called: **the degrees of freedom**.

(more about this tricky concept later).

## Take home message on Sampling Distributions

- Certain *sample statistics* estimate *population parameters*.
  - for example:  $\bar{X}$  estimates  $\mu$ .
- Those statistics are *random variables*, because *every time you sample you get a different outcome!*
- Probability theory tells us the *distribution* of these random variables.
  - for example:  $\bar{X} \sim \mathcal{N}\left(\mu_x, \frac{\sigma^2}{n}\right)$
- This allows us to *infer* something sophisticated about the *population!*

## Take home message on Sampling Distributions

- Certain *sample statistics* estimate *population parameters*.
  - for example:  $\bar{X}$  estimates  $\mu$ .
- Those statistics are *random variables*, because *every time you sample you get a different outcome!*
- Probability theory tells us the *distribution* of these random variables.
  - for example:  $\bar{X} \sim \mathcal{N}\left(\mu_x, \frac{\sigma^2}{n}\right)$
- This allows us to *infer* something sophisticated about the *population!*

## Take home message on Sampling Distributions

- Certain *sample statistics* estimate *population parameters*.
  - for example:  $\bar{X}$  estimates  $\mu$ .
- Those statistics are *random variables*, because *every time you sample you get a different outcome!*
- Probability theory tells us the *distribution* of these random variables.
  - for example:  $\bar{X} \sim \mathcal{N}\left(\mu_x, \frac{\sigma^2}{n}\right)$
- This allows us to *infer* something sophisticated about the *population!*

## Take home message on Sampling Distributions

- Certain *sample statistics* estimate *population parameters*.
  - for example:  $\bar{X}$  estimates  $\mu$ .
- Those statistics are *random variables*, because *every time you sample you get a different outcome!*
- Probability theory tells us the *distribution* of these random variables.
  - for example:  $\bar{X} \sim \mathcal{N}\left(\mu_x, \frac{\sigma^2}{n}\right)$
- This allows us to *infer* something sophisticated about the *population!*

## Inference example: Shaq



Do we *know* that Shaq's probability of getting a free throw is 37.4%?

No! We estimate it based on him making 46 shots out of 123 tries.

$$\begin{aligned}\hat{p} &= 46/123 = 0.374 \\ &= \frac{1}{n} \sum_{i=1}^n X_i\end{aligned}$$

But is it possible that he gets 46/123 shots with a true  $p = 0.5$  (for example)?

## Inference example: Shaq



Do we *know* that Shaq's probability of getting a free throw is 37.4%?

No! We estimate it based on him making 46 shots out of 123 tries.

$$\begin{aligned}\hat{p} &= 46/123 = 0.374 \\ &= \frac{1}{n} \sum_{i=1}^n X_i\end{aligned}$$

But is it possible that he gets 46/123 shots with a true  $p = 0.5$  (for example)?



## Inference example: Shaq



Do we *know* that Shaq's probability of getting a free throw is 37.4%?

No! We estimate it based on him making 46 shots out of 123 tries.


$$\begin{aligned}\hat{p} &= 46/123 = 0.374 \\ &= \frac{1}{n} \sum_{i=1}^n X_i\end{aligned}$$

But is it possible that he gets 46/123 shots with a true  $p = 0.5$  (for example)?


# Possible probability mass functions of 123 baskets

Is it possible that he gets 46/123 shots with a true  $p = 0.5$ ?

**123 free throws at  $p=0.374$**



**123 free throws at  $p=0.50$**



Yes! (Almost) anything is *possible*.

### Question

- How good is  $\hat{p} = \bar{X}$  as an **estimator** of the true  $p$ ?

### A more “formal” statement

- What is the **sampling distribution** of  $\hat{p} = \bar{X}$

### Question

- How good is  $\hat{p} = \bar{X}$  as an **estimator** of the true  $p$ ?

### A more “formal” statement

- What is the **sampling distribution** of  $\hat{p} = \bar{X}$

## Expectation of $\hat{p}$

Assumption:  $X \sim \text{Bernoulli}(p)$  and  $X$  are i.i.d. (independent and identically distributed).

$$\begin{aligned} E(\hat{p}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) \\ &= \frac{1}{n} \sum_{i=1}^n p \\ &= \frac{1}{n} np = p \end{aligned}$$

## Expectation of $\hat{p}$

Assumption:  $X \sim \text{Bernoulli}(p)$  and  $X$  are i.i.d. (independent and identically distributed).

$$\begin{aligned} E(\hat{p}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) \\ &= \frac{1}{n} \sum_{i=1}^n p \\ &= \frac{1}{n} np = p \end{aligned}$$

## Variance of $\hat{p}$

Assumption:  $X \sim \text{Bernoulli}(p)$  and  $X$  are i.i.d. (independent and identically distributed).

$$\begin{aligned}\text{Var}(\hat{p}) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\&= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\&= \frac{n}{n^2} \text{Var}(X_i) \\&= \frac{p(1-p)}{n}\end{aligned}$$

## Variance of $\hat{p}$

Assumption:  $X \sim \text{Bernoulli}(p)$  and  $X$  are i.i.d. (independent and identically distributed).

$$\begin{aligned}\text{Var}(\hat{p}) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\&= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\&= \frac{n}{n^2} \text{Var}(X_i) \\&= \frac{p(1-p)}{n}\end{aligned}$$



## Let's consider some examples:

- Shaq shoots two free throws, and makes 1 of 2.

$$\hat{p} = (0 + 1)/2$$

$$\text{Var}(\hat{p}) = (1/2)(1/2)(1/2) = 0.125$$

$$\text{s.d.}(\hat{p}) = \sqrt{0.125} = 0.3535$$

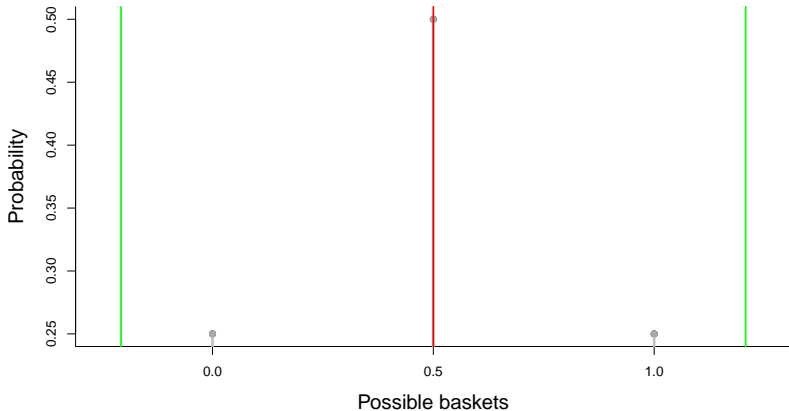
So the estimate of  $\hat{p} = 0.5 \pm 0.707$  is not very precise.

### Important note

- People often report estimates as  $\hat{\mu} \pm 2\hat{\sigma}$
- The range:  $(\hat{\mu} - 1.96\hat{\sigma}, \hat{\mu} + 1.96\hat{\sigma})$  is referred to as: the **95% Confidence Interval**

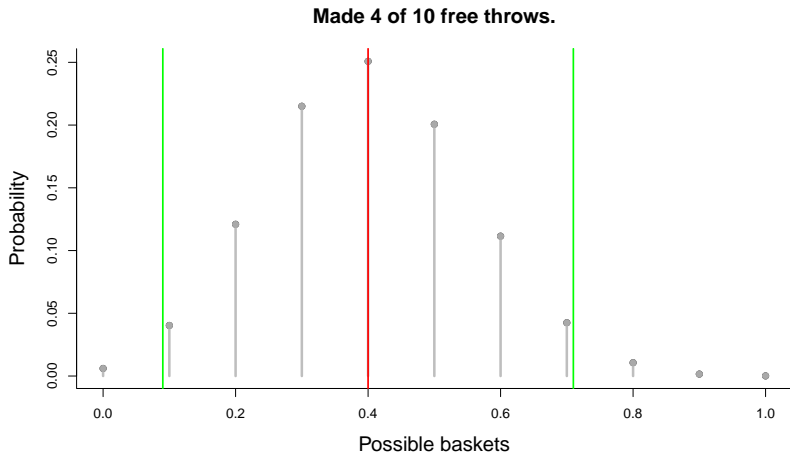
Let's consider some examples:

**Made 1 of 2 free throws.**




$$\hat{p} = 0.5 \pm 0.707, \text{ C.I.} = (-0.21, 1.21)$$

Let's consider some examples:



$$\hat{p} = 0.40 \pm 0.31, \text{ C.I.} = (0.09, 0.71)$$


Let's consider some examples:



$$\hat{p} = 0.35 \pm 0.21, \text{ C.I.} = (0.14, 0.56)$$

Let's consider some examples:


**Made 46 of 123 free throws.**



$$\hat{p} = 0.374 \pm 0.087, \text{ C.I.} = (0.29, 0.46)$$

## Inference question


Given that Shaq shot 46/123 free throws, could he have been a 50% free throw shooter?



- Parameter estimation:  $\hat{p} = 0.374 \pm 0.087$ , C.I. = (0.29, 0.46).
- 50% is outside of the **Confidence Interval** of our estimate ... so very, very unlikely (< 5% chance).

## Inference question


Given that Shaq shot 46/123 free throws, could he have been a 50% free throw shooter?



- Parameter estimation:  $\hat{p} = 0.374 \pm 0.087$ , C.I. = (0.29, 0.46).
- 50% is outside of the **Confidence Interval** of our estimate ... so very, very unlikely (< 5% chance).

## Inference question

Given that Shaq shot 46/123 free throws, could he have been a 50% free throw shooter?



- Parameter estimation:  $\hat{p} = 0.374 \pm 0.087$ , C.I. = (0.29, 0.46).
- 50% is outside of the **Confidence Interval** of our estimate ... so very, very unlikely (< 5% chance).